

FlowQoS: Not Every Flow is Born the Same

Dhruv Sharma
UC San Diego, CA
dhsharma@cs.ucsd.edu

Robert Jenkins
UC San Diego, CA
rfjenkins@ucsd.edu

Frederik Nygaard
UC San Diego, CA
frederikny@gmail.com

Feichao Qian
UC San Diego, CA
feqian@ucsd.edu

Abstract—In modern times, user devices connected to broadband access networks, run an assortment of applications that exchange network traffic with remote servers and other devices over the Internet. However, in the absence of any differentiated treatment by the network devices, traffic from certain delay and loss-sensitive applications suffers badly leading to performance degradation for such applications. In this paper, we validated the classification and traffic isolation technique proposed by the original FlowQoS paper and conclude that it certainly improves the bitrates as well as reduces packet jitter and latencies for certain video streaming applications that provide the end-user with a much better experience. Also, we propose and evaluate an alternate architecture for FlowQoS using Linux traffic control disciplines and evaluate it demonstrate the additional benefits that can be gained by adopting such approach.

I. INTRODUCTION

In modern times, user devices connected to broadband access networks, run an assortment of applications that exchange network traffic with remote servers and other devices over the Internet. Since the upstream and downstream throughput is generally limited, these traffics will compete for relatively scarce bandwidth resources.

However, traffic from one application might not share the same characteristics as the traffic created from another application. To a large extent, it is the end-user's requirements that play a significant role in deciding the nature of such applications and hence the kind of network traffic they send and receive. For instance, a user's expectation from a VoIP call, video streaming and gaming applications is that their experience remains seamless and of high quality during their use of the application, requiring the traffic to be sent out at near constant rates with low delay and high reliability paths over the network. Whereas, in certain use cases such as data backup to the cloud and system updates, it is expected that the operation must completed eventually - even when the user is not actively using the application. Traffic from such applications, in contrast to VoIP and video streaming, does not come with any hard deadlines or network requirements.

While the nature of traffic varies so widely, the network devices today are, to a large extent, agnostic to such subtleties in the nature of network traffic and tend to handle it in the same way. Even though later kinds of applications discussed above do not face any issues that directly affect the user, the former kind might be impacted severely due to the effect on network dynamics leaving the user with a sour experience.

Research over the past years has resulted in the use of various metrics such as packet delays, jitter, available bandwidth,

frame rate etc. to quantify users experience in some way. One possible way to deal with limited throughput is to configure the network routers to prioritize some specific applications' traffic flows (e.g. video, VoIP etc) over others (e.g. data backup, file upload etc). It will effectively improve Quality of Service (QoS). However, for some reason, previous work on QoS mechanisms have not been deployed in broadband access networks [22]. The emergence of software-defined networking (SDN) gives more possibilities to solve this problem. One approach to deploying QoS in broadband access networks is to utilize the advantages of SDN that separates the network's control logic and forwarding planes. We can migrate the functions that perform QoS both application identification and router-level configuration to separate control logic, and design a front-end client (e.g., webpage) at high levels of abstraction to let users to assign bandwidth to each identified application according to their own preference. Once users set up their preference, the front-end will install the QoS configuration into the home routers [22]. This approach makes it easy for user to configure priorities and facilitates more sophisticated per-flow application based QoS.

A. Our Contribution

First, we deploy the FlowQoS implementation of pair of virtual switches for improving QoS using classification and traffic policing, within a virtual machine (emulating our end-host) with Internet connection instead of the discussed hardware implementation in the original FLOWQoS paper. Second, we validate the results presented by FlowQoS [22] under similar scenarios that they used but evaluating the effects using different metrics and applications. Finally, we implement and evaluate an architecture that utilizes Linux's advanced routing and traffic control to implement the idea of FlowQoS while overcoming the limitation of under-utilization of available bandwidth.

B. Outline

The rest of the paper is organized in the following way: We discussed related work in QoS, as well as SDN-based solutions for home and broadband access networks in § II. In § III we discuss the difference between rate limiting using traffic policing and traffic shaping. The motivation of FlowQoS is presented in § IV. § V describes various components of the FlowQoS implementation. In § VI, we evaluate FlowQoS for video streaming and VoIP applications in the context of competing TCP flows. Finally, we discuss our results, future

work and other open research avenues in § VII and conclude in § VIII.

II. RELATED WORKS

Quality of Service (QoS) is important in home networks that share a common bandwidth-limited Internet connection. Over the past several years, there has been a significant amount of research for QoS in IP networks [11], [16], [19], traffic shapers [5], and traffic flows classifiers [21]. However, most previous approaches are different with FlowQoS either focusing on different issues or working in different scenarios. FlowQoS focuses in particular on making per-flow, application-based QoS, which is designed to deploy and configure in home networks.

Kim et al. provides a network QoS control framework [15], which sets rate limiters at the edge switches and priority queues for flow at each path hop. It uses a QoS control framework to manage automatically OpenFlow networks with multiple switches. On the contrary, FlowQoS provides similar automated traffic shaping at a single gateway. Ishimori et al. introduces a novel QoSFlow framework to enhance QoS management procedures in OpenFlow networks [14]. QoSFlow shares similar shaping mechanism with FlowQoS. However, it does not focus on providing usable QoS for broadband access networks and is still under development. Ko et al. proposed a two-tier flow-based QoS management framework [18], which needs multicore processors and is not designed for home networks like FlowQoS. Ferguson et al. developed an API for applications to control a software-defined network (SDN). They call their prototype controller PANE [12]. PANE delegates read and write authority from the networks administrators to end users, or applications and devices acting on their behalf. It allows users to work with the network, rather than around it. However, PANE addresses on more issues than FlowQoS such as better performance, security and fairness etc., which makes it do not focus on QoS in broadband access networks or application identification. Williams et al. [23] developed an automated IP traffic classification algorithm based on statistical flow properties. This approach limits the throughput of commodity home routers to 28 Mbps. In contrast, FlowQoS faces no such limitations.

Current residential router QoS support must be manually configured and set up is difficult for the average home users. The emergence of OpenFlow/SDN is one of the most promising and disruptive networking technologies of recent years. It provides more possible solutions for QoS. In the SDN architecture, the control and data planes are decoupled. The underlying network infrastructure is abstracted from the application, which mitigate the issue that most users are not skilled network operators. There is also a lot of work based on SDN. Risso et al. [20] developed an OpenFlow-based mechanism for customizing data-plane processing in home routers, which makes it possible to install third parties' applications on the data plane of a router. The architecture is focused on more general data-plane modifications, not QoS. Georgopoulos et al. [13] proposed an OpenFlow-assisted

QoE fairness framework that aims to fairly improves users quality of experience (QoE) of multiple competing clients in a shared environment like home networks. However, the system performs per-device QoS instead of per-application or per-flow QoS. Mortier et al. [17] designed and built a home networking platform named Homework to provide detailed per-flow measurement and management capabilities for home networks. However, Homework does not provide QoS support or application identification.

III. BACKGROUND

Before we discuss FlowQoS's overall design we would like to introduce the concept of traffic policing and traffic shaping. While both techniques can be used for effectively limiting the rate at which traffic is forwarded through a link, these fundamentally differ in the way they handle the network packets.

Traffic policing is usually implemented using token-based approach where based on the set rate limit, a fixed number of tokens are allocated to an interface at the end of each token refresh period. As packets of corresponding sizes are forwarded the token count is reduced until no more tokens can be removed. This ensures that the interface does not send traffic more than the stipulated rate. However, if there are still packets left with the interface that could not be sent, the policer simply drops those packets. Also due to token based nature of policer, packet bursts get forwarded by the interface. The OpenvSwitch (OVS) [6] implements rate limiting using traffic policing that is used by the FlowQoS's original architecture.

Traffic shaping, on the other hand, is implemented using a leaky-bucket approach (one variant of this is equivalent to a token-based approach) which in fact calculate an equivalent rate metric based on system architecture timer resolution and forwards traffic continuously strictly conforming to this rate. Any burst of packets (packets at higher rate) are buffered in a short internal queue and delayed until the shaper achieves the desired transmission rate. Only when the queue becomes full, the residual packets are dropped. Due to this queuing, packets are expected slightly higher latencies than normal. The Hierarchical Token Bucket (HTB) [3] queuing discipline provided by the Linux's traffic control suite implements traffic shaping that follows a similar approach to forward traffic.

Moreover, HTB is classful queuing discipline that allows adding multiple traffic filter (classes) under the root filter (class), which can further have sub-filters (sub-classes) forming a traffic classification tree. Each of the intermediate and leaf classes can have individual rate limits and therefore traffic shapers associated with them. In addition, they also have an associated priority that is respected by HTB while sharing unutilized bandwidth among multiple sibling classes. Fig. 1 illustrates such a hierarchical structure described above and shows how egress traffic is partitioned into multiple classes.

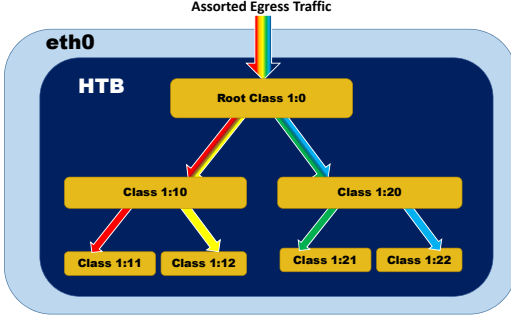


Fig. 1: HTB Hierarchical Class Structure Illustration.

IV. MOTIVATION

Congestion has been, and continues to be, one of the most prevalent problems in networking. One of the key issues with TCP is that when there are many flows it has a tendency to cause congestion, leading to degraded performance of competing flows. This degradation can range in cost from minimal to disastrous. When TCP connections consume all the available bandwidth real-time applications have a tendency to suffer the most. This is true because real-time applications have tighter latency and throughput constraints.

Our motivation for exploring FlowQoS is to provide better quality streaming services, which cannot afford to compete with flows like TCP which can dominate them. It is commonplace for many users to have multiple web pages with active web traffic while simultaneously using a streaming service like Skype or VoIP- we believe that flow classification described in FlowQoS can prevent these users from receiving poor Skype call quality due to congestion that they may be creating.

One of the downsides of FlowQoS is that it statically allocates bandwidth to classes of traffic, such as TCP or UDP. This is detrimental to throughput when certain classes are not fully utilizing their allotted bandwidth. By using the concept of classifying traffic from FlowQoS and dynamically allocating bandwidth we will show that we can still achieve get the target quality of service that FlowQoS provides, while providing a total throughput that is closer to a non-regulated network.

V. DESIGN

The original FlowQoS prototype was implemented on Raspberry Pi running OpenWrt Linux distribution with Openvswitch integration. However, since the aim of this paper is to validate the claims and results that demonstrate the need for FlowQoS-like systems rather than its deployment feasibility, we implement our solution within the virtual networking environment in end-host Linux distribution. Though such implementation can be still be easily moved to a dedicated device.

The FlowQoS's OVS-based architecture (Fig. 2) is comprised of dual-OVS topology with multiple paths between

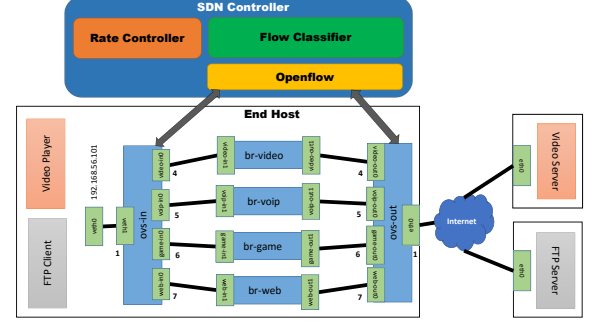


Fig. 2: FlowQoS architecture.

them passing through "dumb" forwarding bridges. Each path between the two OVSes (ovs-in and ovs-out) corresponds to a high-level traffic class to which the user or administrator wishes to classify any ingress (or egress) traffic to (or from) the end-host. Both ovs-in and ovs-out are configured to run in secure Openflow mode such that the OVS would accept Openflow commands from the Openflow controller when connected to it and in its absence it would not default to learning switch mode. This is required to prevent the runaway issue in learning switch network topologies with cycles. The physical interface (eth0 in our case) connecting the end-host to the physical network is attached to ovs-out, whereas, the IP address for the end-host machine is moved from eth0 to the virtual ethernet interface veth0. Such a setup ensures that any traffic from any application must traverse dual-OVS virtual topology before leaving the end-host and vice-versa for the incoming traffic. While the connecting bridges between ovs-in and ovs-out provide no functionality of their own, we use them solely for packet capture and debugging purposes.

In our implementation of HTB-based architecture, we combined ovs-in and ovs-out in Fig. 2 into one OVS and enable an HTB qdisc (queueing discipline) on veth0 and veth1. To this qdisc, we added one parent class and four leaf classes each for one traffic type. The leaf classes are: **Video**, **VoIP**, **Gaming** and **Web**. To this, we added filters to classify the different traffic types and finally assign priorities to the each sub-class such that **Video** and **VoIP** get the highest priorities followed by **Web** traffic and finally the **Gaming** (unclassified) traffic gets the least priority.

We now discuss the two primary components of the FlowQoS controller module - namely, the Flow Classifier and the Rate Controller.

A. Flow Classifier

FlowQoS's classifier is comprised of two sub-modules - HTTP classifier and non-HTTP classifier. The reason for such initial segregation is that modern HTTP applications no longer just carry textual web traffic but can include embedded components such as video players, advertisements, images

etc. each generating traffic of different nature. Therefore, identifying traffic to be HTTP is not sufficient and further analysis must be done to ascertain the type of traffic that the HTTP stream is actually carrying. Hence, any incoming packets from underlying Openflow layer are categorized HTTP (or HTTPS) based on the source or destination port being 80 (or 443).

With this goal in mind, FlowQoS was designed to inspect DNS queries that an HTTP application makes before it even sends or receives any HTTP traffic. The classifier uses a preconfigured list of DNS CNAME records to match a domain to its application type. The list is continuously updated based on the expiry time of DNS record. Hence, when an HTTP traffic has source or destination that matches the regular expression for that domain, it is classified based on the pre-identified traffic type.

On the other hand, for non-HTTP traffic, FlowQoS makes use of the `libprotoident` [2] library. The library requires certain additional information to be captured from the packets apart from the protocol type, source destination IP addresses and ports. Particularly, it uses the first 4 payload bytes and payload size for the first packet in either direction flows of a connection (or just one direction in case of unidirectional flows).

However, our initial experimentation with applications we use to evaluate FlowQoS did not result in correct identification of the protocol due to multiple reasons such as unavailability of DNS resolution for locally hosted HTTP server and due to certain protocols (eg. Skype and Iperf TCP) not showing specific packet signature for classification by `libprotoident` library. Therefore, for our evaluation we fall back to the default port-based protocol identification technique.

B. Rate Controller

FlowQoS's rate controller hosts REST APIs to receive a bandwidth distribution from the user and generate a parsable JSON configuration file to be used on the end-host to enforce rate limits on each kind of traffic. Since, the primary focus of this paper is not to validate the accuracy of the flow classification, but rather the effectiveness of classification and rate limiting in improving application performance, we now evaluate the two different schemes of enforcing rate limits along with their pros and cons based on the evaluations with various traffic types in § VI.

VI. EVALUATION

In this section, we evaluate the effectiveness of multi-channel OpenvSwitch design in improving the performances of namely three audio-video streaming applications. Each of these applications produce delay and loss-sensitive video traffic. However, the sensitivity to these factors varies differently for these applications and hence the metrics to measure user-satisfaction are also different for each of these applications.

In addition, we evaluate and discuss the improvements that our proposed hierarchical token bucket (HTB) design can bring

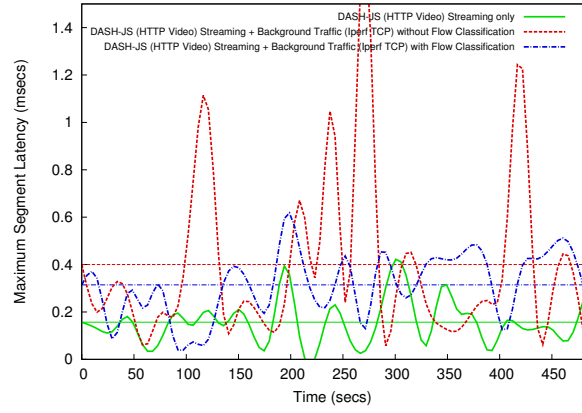


Fig. 3: Reduction in maximum segment latencies with DASH streaming using FlowQoS's OVS-based architecture

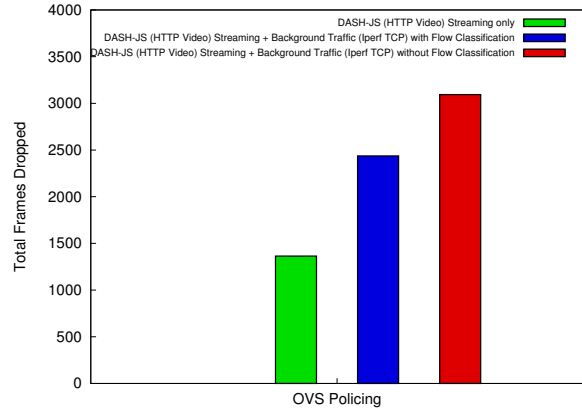


Fig. 4: Reduction in number of frames dropped with DASH streaming using FlowQoS's OVS-based architecture

by replacing the OVS-based scheme while performing the similar classification and rate-limiting functionality as before.

Now, we discuss the three applications that we used to for our evaluation and discuss their results. For each application we measure the corresponding performance metrics under the following three scenarios:-

- The application alone sends (or receives) traffic and gets maximum available bandwidth.
- The application competes for bandwidth with a heavy, long-running TCP-based background traffic in the absence of FlowQoS scheme.
- The application and the heavy, long-running TCP-based background traffic share 50% of the maximum available bandwidth each by employing FlowQoS's classification and rate-limiting scheme.

To emulate a real-life TCP-based background traffic we used the Iperf utility [4] available in all major Linux distributions. The utility is used to measure TCP or UDP throughput and can run in both client and server mode depending on the command-line arguments. As a server it accepts connections at specified port and receives high amounts traffic on that connection, whereas as a client it connects to specified server address and sends high amount of TCP or specified amount of UDP traffic

to the Iperf server.

A. FlowQoS with OVS-based architecture

1) *DASH-JS Player*: The Dynamic Adaptive Streaming over HTTP (DASH) [1] is an adaptive bitrate-based video streaming technique where a large video file is partitioned into small segments worth few seconds (or minutes) of playback time and each segment is encoded in multiple bitrates and stored at the streaming server. The segments are sent as HTTP responses to the video player in end-host's web browser where it is decoded and played. In conditions of congestion or packet loss, DASH-JS player [1] reduces the bitrate requirements and the streaming server starts streaming the subsequent segments at a lower bitrate.

While the original paper, discusses the effectiveness of their scheme on adaptive bitrate and throughput, we evaluate two other important metrics critical to user experience - the segment arrival latency and number of frames dropped. High segment arrival latencies tend to cause low buffer lengths at the video player leading to the video playback freezing frequently. Whereas, the number of dropped frames indicate that even though some initially dropped packets were re-transmitted by underlying TCP layer in HTTP, they did not reach the buffer in time for playback and the were dropped. For a user's perspective, this causes the video to suddenly jump ahead (sometimes giving the experience popularly known as "robotic movements") during playback.

The results were collected by setting up two VMs running Ubuntu. One of the VMs acted as a client requesting the video data and hosting an iperf server. The other VM acted as the server, sending the video data as well as running 16 iperf processes to fill up the bandwidth and cause network congestion. Figure 3 shows the max latencies of the segments as they are being streamed from the server VM to the end-host VM. Streaming without flow classification has the largest range of values, and indeed exhibited the worst experience for watching the video. After adding the flow classifier, not only did the streaming perform better, but there was significantly less variation- almost comparable to having no traffic.

The best measure of how these results affect user experience is perhaps in the number of dropped frames. Figure 4 shows that flow classification with background traffic is able to outperform the same scenario without flow classification. This result corroborates the work from the FlowQoS paper that flow classification can be used to prioritize traffic types to give the best user-experience.

2) *VLC Real-time Streaming*: The VLC media player [9] is a portable, free and open-source cross-platform media player and streaming media server developed as a part of the VideoLAN [9] project. For our evaluation we used a VLC server-client player to stream the Big-Buck-Bunny video (often used to compare multiple video standards, encoding schemes and protocols) over a private network using the Real-time Transport Protocol (RTP) [7] with MPEG encoding. Unlike DASH, VLC's video streaming server uses UDP as transport protocol for RTP, sends data in single encoding

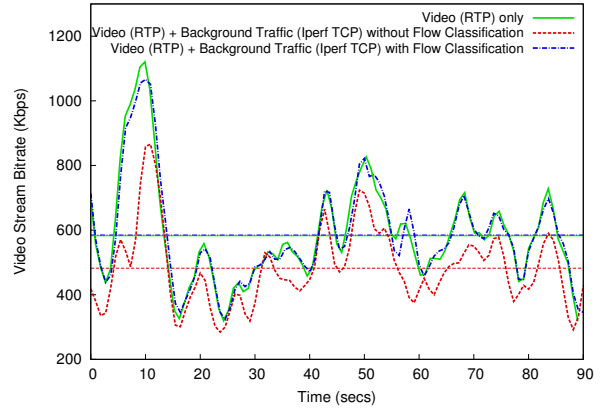


Fig. 5: Improved VLC (RTP) video streaming bitrate with FlowQoS's OVS-based architecture.

format and maintains a small playback buffer. Therefore, during congestion, it suffers higher frame drops compared to DASH due to UDP packet loss and small playback buffer size. Also, due to single encoding format, the server cannot vary the bitrate to adjust reduce the packet loss suffered during the streaming leading to poor quality video for its complete playback time.

For evaluating the impact of FlowQoS on the RTP-based stream performance, our setup was similar to the experiment for DASH-JS player. We had two VMs, one for end-host running VLC RTP client and Iperf TCP server, and the other VM running VLC RTP server and Iperf TCP client such that we have high amount of both RTP and regular TCP traffic coming in to the end-host VM. Fig. 5 shows bitrate variation for the three scenarios we described in § VI. Clearly, bitrate for RTP stream is lower when it has to compete with the TCP stream for bandwidth. However, FlowQoS based traffic isolation actually helps us achieve a bitrate values equivalent to the case where RTP has no competing traffic. While it only seems to be a 100Kbps of bitrate difference, the video quality from a user's perspective is incredibly poor when FlowQoS is not used for flow isolation leading to large sections of video screen being blacked-out.

3) *Skype*: Skype [8] is a popular Voice over IP (VoIP) service, recently acquired by Microsoft, that provides real-time audio, video calling and text messaging functionality using a closed source proprietary protocol running over the network (IP) layer. Currently, it forwards the video (and audio) traffic between two clients via an hybrid peer-to-peer overlay of Microsoft super-nodes spread across the globe.

Compared to DASH and RTP-based video streaming, the quality and real-time constraints on Skype video calls are much harder requiring high stability in the birates and packet interarrival times to keep interaction between two individuals as seamless as possible. Therefore, we take the variation in packet interarrival times at the client side as a measure of performance for FlowQoS implementation. This measure is often known as "jitter" and is measured in units of time (usually milliseconds). Skype clients often come with a feature

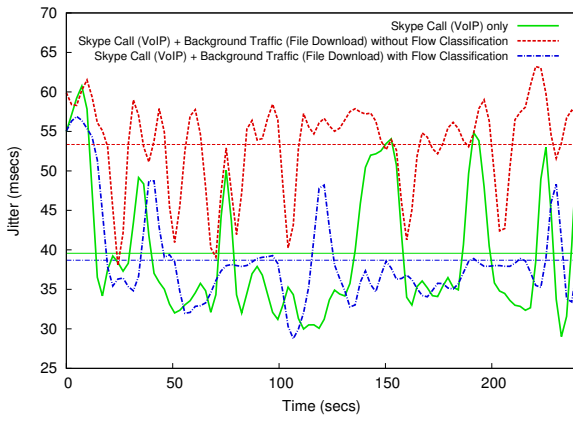


Fig. 6: Reduced Skype Video Call Jitter with FlowQoS's OVS-based architecture.

to enable real-time in-call statistics reporting jitter, round-trip time (RTT), capture frame rate, bitrate etc. We use the jitter values reported by Skype's in-call statistics.

However, to ensure that minor network variations do not vary jitter values significantly, we chose to conduct a long distance international video call that ensured relatively stable jitter and RTT values with Skype traffic alone, and we simultaneously monitored for changes in the relay path that could affect these measures.

Unlike, the cases with DASH-JS player and VLC video streaming, Skype requires Internet access. Hence, for this experiment we used one VM as the end-host running Skype client as well as a *wget* [10] utility as the FTP client downloading a large OS image from a remote FTP server. As shown in Fig. 6, we observed that both Skype's UDP traffic competes with the incoming FTP (TCP) traffic, the jitter for the Skype call is not only high on an average but it also varies drastically over time. On the other hand, using FlowQoS's OVS-based isolation scheme we could achieve lower and more stable jitter values. The result show with FlowQoS we can achieve even lower jitter compared to the case where Skype is run standalone. We believe such a result can be attributed to the variation in network conditions as well as effect of queuing in FlowQoS OVSes, though further investigation might be required to validate this claim.

B. FlowQoS with HTB-based Architecture

1) *Efficient Bandwidth Utilization:* While, the results with the FlowQoS's OVS-based architecture were impressive, it still suffers from a significant drawback of under-utilizing the network bandwidth in the absence of any competing higher priority traffic. Fig. 7 shows such a scenario. From 0-40 secs interval, video traffic gets the complete bandwidth for it's use. Between 40-80 secs interval, a competing Iperf TCP traffic runs in parallel getting 50% of the bandwidth share. However, during the 80-120 secs interval, the video traffic (high priority) has stopped but the TCP still gets 50% bandwidth share leading to a 50% wastage of available bandwidth.

In comparison, by using a Hierarchical Token Bucket approach, the ceiling rate and priority values for each HTB

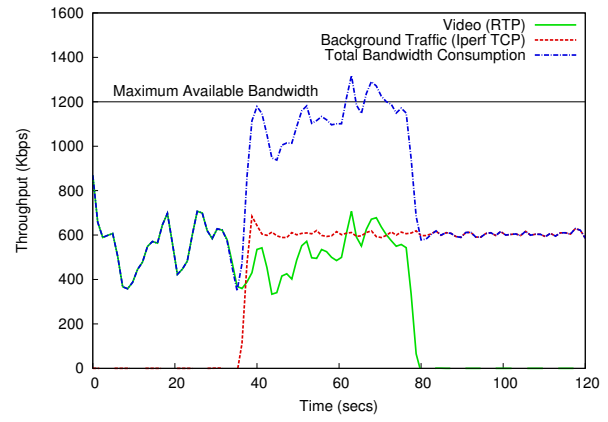


Fig. 7: Inefficient Bandwidth Utilization with FlowQoS's OVS-based architecture.

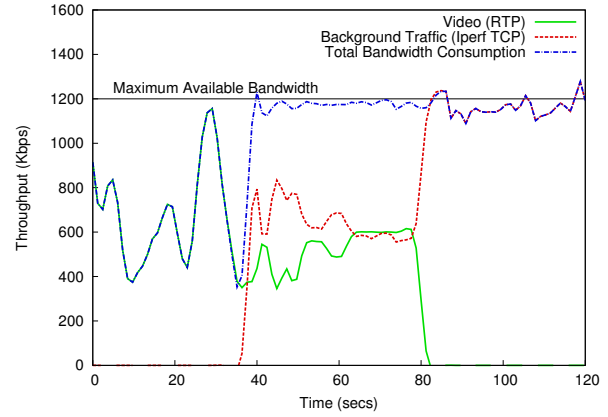


Fig. 8: Efficient Bandwidth Utilization with FlowQoS's HTB-based architecture.

traffic can be configured to enable a high priority class to share it's unused bandwidth with an immediate lower priority class. Fig. 8 shows a similar scenario we saw earlier, however, we can clearly observe that during 80-120 secs interval, the absence of video traffic gives the total available bandwidth to the lower priority TCP traffic. This leads to an efficient utilization of available bandwidth compared to the OVS-based solution.

2) *DASH-JS Video Player:* We re-experimented with the DASH-JS video player on a similar server-client setup as before, but used FlowQoS's HTB-based architecture instead. We observed that not only the average segment latencies decreased for the scenarios in which video traffic was isolated but also observed that the overall variation (measured by the standard deviation) in the latencies also reduced. This can be attributed to the fact that in either case video traffic does not saturate its class's bandwidth share and therefore, HTB did not require to buffer any packets and forwarded a continuous stream immediately, whereas due to the relative simplicity of the OVS rate-limiting implementation, it is less accurate in rate estimation and handling fragmented packets leading to higher latencies. Also, we see that for the case where both video traffic compete, HTB architecture gives higher latencies. We

VII. DISCUSSION & FUTURE WORK

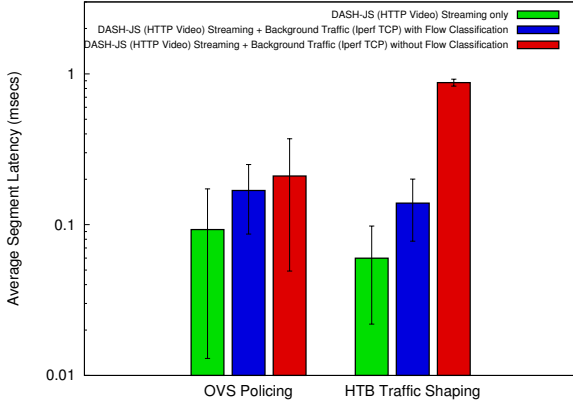


Fig. 9: Reduction in mean latencies for classified flows using FlowQoS's HTB architecture

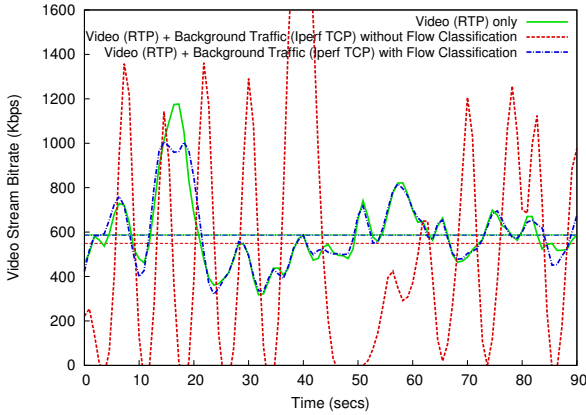


Fig. 10: Improved overall bitrate for unclassified flows using FlowQoS's HTB architecture

attribute this to the fact that such a situation leads to congestion on that link causing buffering in HTB's traffic shaper that adds queuing delays to the packets. Whereas in case of OVS's policing, such packets are dropped and since DASH-JS player only reports latencies for segments received, OVS seems to perform better even though it would see higher packet loss rate.

3) *VLC Video Streaming*: We repeated the three scenarios mentioned in § VI with the FlowQoS's HTB-based architecture in place for VLC-based video streaming experiment as shown in Fig. 10. We observed that HTB-based approach indeed produces bitrates equivalent to the OVS-based architecture when we classify and isolate traffic along multiple paths. Moreover, it also improves the overall bitrate of the VLC video streaming even when we do not isolate it from TCP traffic. This can be attributed to the traffic shaping nature of HTB due to which it prefers delaying packets rather than dropping them leading to a higher overall bitrate. Eventhough, this leads to high variation in the bitrate over time, we observed that due an overall increase in mean bitrate user experiences a relatively good video quality.

Through our experiments, we observed that the FlowQoS architecture in either case helps in increasing the overall bitrate for VLC-based RTP video streaming, reduces overall jitter for Skype video calls and reduces the overall segment latencies in comparison to the scenario where both of must compete for their share of bandwidths. Moreover, we observed that our alternate implementation of the FlowQoS architecture using HTB queuing routine certainly improves bitrate for video traffic when it competes with a simultaneous TCP traffic, reduces latencies for DASH-JS segment delivery and provides a much more efficient utilization of available bandwidth by a lower priority traffic in case a higher priority traffic is not consuming it's share of bandwidth.

We have shown that classifying and separating flows is an effective way to make sure congestion doesn't ruin performance requirements for applications that are more affected by longer latencies.

One major concern seen by our results is the case where flows shared the same queue as the TCP traffic. In the Dash-JS experiment we saw that the HTB shows significantly higher latencies when flow remain unclassified and share the same link when compared to the OVS policing case. We defer it for a future investigation of this discrepancy and answer whether (or not) the additional overhead introduced by it could significantly affect its scalability compared to the OVS-based solution.

In the future, we would like to exploit the NetFlow support provided in OpenvSwitch for realtime monitoring, analytics and dynamic rate control for various traffic types using a local end-host agent. This can help in efficiently utilizing the available bandwidth in OVS-based solution itself similar to the HTB-based solution.

Moreover, the current solution does not consider equal bandwidth share among flows within each traffic type and this remains an open problem for us. We intend to explore the Stochastic Fair Queuing (SFQ) discipline in Linux's traffic control suite to enforce this constraint.

Also, we aim to experiment with other protocol identification tools libraries that use more sophisticated classification techniques and achieve higher accuracy in identification of protocols compared to `libprotoident`.

VIII. CONCLUSION

In this paper, we validated the results obtained by the original FlowQoS paper [22] that demonstrate the isolation of flow traffic based on some high-level classification and bring significant improvement to the performance and user experience for certain packet loss and delay sensitive real-time applications. In addition, we propose an alternate architecture using Linux's traffic control suite's Hierarchical Token Bucket and demonstrate that in general not only it performs at least as well as it's OVS-based counterpart, it can also improve the overall utilization of available bandwidth for the end-host.

REFERENCES

- [1] Dash. <http://dashif.org/>.
- [2] FlowQoS. <http://flowqos.noise.gatech.edu/>.
- [3] HTB Linux queuing discipline manual - user guide. <http://luxik.cdi.cz/~devik/qos/htb/manual/userg.htm>.
- [4] Iperf utility. <http://manpages.ubuntu.com/manpages/trusty/man1/iperf.1.html>.
- [5] Meraki Traffic Shaping Technology. https://meraki.cisco.com/lib/pdf/meraki_whitepaper_layer_7_visibility_and_control.pdf.
- [6] OpenvSwitch. <http://openvswitch.org/>.
- [7] RTP - A Transport Protocol for Real-Time Applications. <https://www.ietf.org/rfc/rfc3550.txt>.
- [8] Skype. <http://skype.com/>.
- [9] VLC Median Player and Server. <http://www.videolan.org/>.
- [10] Wget utility. <http://manpages.ubuntu.com/manpages/trusty/man1/wget.1.html>.
- [11] C. Aurrecochea, A. T. Campbell, and L. Hauw. A survey of qos architectures. *Multimedia systems*, 6(3):138–151, 1998.
- [12] A. D. Ferguson, A. Guha, C. Liang, R. Fonseca, and S. Krishnamurthi. Participatory networking: An api for application control of sdns. In *ACM SIGCOMM Computer Communication Review*, volume 43, pages 327–338. ACM, 2013.
- [13] P. Georgopoulos, Y. Elkhatib, M. Broadbent, M. Mu, and N. Race. Towards network-wide qoe fairness using openflow-assisted adaptive video streaming. In *Proceedings of the 2013 ACM SIGCOMM workshop on Future human-centric multimedia networking*, pages 15–20. ACM, 2013.
- [14] A. Ishimori, F. Farias, I. Furtado, E. Cerqueira, and A. Abelém. Automatic qos management on openflow software-defined networks. *IEEE Journal*, 2012.
- [15] W. Kim, P. Sharma, J. Lee, S. Banerjee, J. Tourrilhes, S.-J. Lee, and P. Yalagandula. Automated and scalable qos control for network convergence. *Proc. INM/WREN*, 10:1–1, 2010.
- [16] D. McDysan. *QoS and traffic management in IP and ATM networks*. McGraw-Hill, Inc., 1999.
- [17] R. Mortier, B. Bedwell, K. Glover, T. Lodge, T. Rodden, C. Rotsos, A. W. Moore, A. Kolioussis, and J. Sventek. Supporting novel home network management interfaces with openflow and nox. In *ACM SIGCOMM Computer Communication Review*, volume 41, pages 464–465. ACM, 2011.
- [18] K. Nam-Seok, H. Hwanjo, P. Jong-Dae, and P. Hong-Shik. Openqflow: Scalable openflow with flow-based qos. *IEICE transactions on communications*, 96(2):479–488, 2013.
- [19] P. Newman, W. Edwards, R. Hinden, E. Hoffman, F. C. Liaw, T. Lyon, and G. Minshall. Ipsilon flow management protocol specification for ipv4 version 1.0. Technical report, 1996.
- [20] F. Risso and I. Cerrato. Customizing data-plane processing in edge routers. In *Software Defined Networking (EWSN), 2012 European Workshop on*, pages 114–120. IEEE, 2012.
- [21] M. Roughan, S. Sen, O. Spatscheck, and N. Duffield. Class-of-service mapping for qos: a statistical signature-based approach to ip traffic classification. In *Proceedings of the 4th ACM SIGCOMM conference on Internet measurement*, pages 135–148. ACM, 2004.
- [22] M. S. Seddiki, M. Shahbaz, S. Donovan, S. Grover, M. Park, N. Feamster, and Y.-q. Song. FlowQoS: Per-Flow Quality of Service for Broadband Access Networks. *Proceedings of the third workshop on Hot topics in software defined networking - HotSDN '14*, pages 207–208, 2014.
- [23] N. Williams and S. Zander. Real time traffic classification and prioritisation on a home router using diffuse. Technical report, CAIA Technical Report 120412A, 2011.